

# **Kernel-Based Learning & Multivariate Modeling**

**MIRI Master - DMKM Master**

**Lluís A. Belanche**  
belanche@cs.upc.edu

Soft Computing Research Group

*Universitat Politècnica de Catalunya*

2016-2017

# Kernel-Based Learning & Multivariate Modeling

## Contents by lecture

**Sep 14** Introduction to Kernel-Based Learning

**Sep 21** The SVM for classification, regression & novelty detection (I)

**Sep 28** The SVM for classification, regression & novelty detection (II)

**Oct 05** Kernel design (I): theoretical issues

**Oct 19** **Kernel design (II): practical issues**

**Oct 26** Kernelizing ML & stats algorithms

**Nov 02** Advanced topics

# Kernel design (II): practical issues

## Euclidean space $\mathbb{R}^d$ , but not only

- Kernels on real vectors (whole families)
- Kernels on binary vectors (bitstrings = sets)
- General structured kernels:
  - All-subsets kernel
  - Convolution kernels
- Kernels on discrete structures:
  - Tree kernels
  - Graph kernels
- Kernels on distributions (generative kernels):
  - P-kernels
  - Marginalized kernels
- String kernels (text)

... and many others (functional data, categorical data, ...)

# Kernel design (II): practical issues

## All-subsets kernel

Consider a feature space with one feature for every subset  $A \subseteq \{1, \dots, d\}$  of the input variables:

for  $\mathbf{x} \in \mathbb{R}^d$ , the feature  $A$  is given by  $\phi_A(\mathbf{x}) = \prod_{i \in A} x_i$  (note  $\phi_\emptyset(\mathbf{x}) = 1$ )

The kernel is defined by the mapping  $\phi : \mathbf{x} \rightarrow (\phi_A(\mathbf{x}))_{A \subseteq \{1, \dots, d\}}$

$$\begin{aligned} k(\mathbf{x}, \mathbf{x}') &= \langle \phi(\mathbf{x}), \phi(\mathbf{x}') \rangle = \sum_{A \subseteq \{1, \dots, d\}} \phi_A(\mathbf{x}) \phi_A(\mathbf{x}') \\ &= \sum_{A \subseteq \{1, \dots, d\}} \prod_{i \in A} x_i x'_i = \prod_{i=1}^d (1 + x_i x'_i) \end{aligned}$$

---

The last step is obtained by expanding  $(1 + x_1 x'_1)(1 + x_2 x'_2) \dots (1 + x_d x'_d)$

# Kernel design (II): practical issues

## All-subsets kernel

We have the freedom to downplay some features (and thus emphasize others) by introducing weighting factors  $w_i \geq 0$  for each feature  $i$ :

$$\phi_A(\mathbf{x}) = \prod_{i \in A} \sqrt{w_i} x_i$$

therefore

$$k_{\mathbf{w}}(\mathbf{x}, \mathbf{x}') = \prod_{i=1}^d (1 + w_i x_i x'_i)$$

# Kernel design (II): practical issues

## Bitstring/Binary variables/Sets

Let  $\mathbf{x}, \mathbf{x}' \in \{0, 1\}^d$ . The *Simple Matching Coefficient* (SMC) is the fraction of 1 – 1 matches, and it is a kernel on  $\{0, 1\}^d$ .

*Proof.* For every  $N \in \mathbb{N}$ , and every choice  $\mathbf{x}_1, \dots, \mathbf{x}_N \in \{0, 1\}^d$ , we form the matrix  $K = (k_{ij})$ , where  $k_{ij} = k(\mathbf{x}_i, \mathbf{x}_j) = \frac{1}{d} \sum_{k=1}^d x_{ik} x_{jk} = \frac{1}{d} \langle \mathbf{x}_i, \mathbf{x}_j \rangle$

Then we have, for any  $\mathbf{c} \in \mathbb{R}^N$ :

$$\sum_{i=1}^N \sum_{j=1}^N c_i c_j \frac{1}{d} \langle \mathbf{x}_i, \mathbf{x}_j \rangle = \frac{1}{d} \left\langle \sum_{i=1}^N c_i \mathbf{x}_i, \sum_{j=1}^N c_j \mathbf{x}_j \right\rangle = \frac{1}{d} \left\| \sum_{i=1}^N c_i \mathbf{x}_i \right\|^2 \geq 0$$

# Kernel design (II): practical issues

## Bitstring/Binary variables/Sets

Given two finite sets  $A, B$ , consider

$$k(A, B) = \sum_{a \in A} \sum_{b \in B} k_{\text{base}}(a, b)$$

If  $k_{\text{base}}$  is the overlap kernel:

$$k(a, b) = \begin{cases} 1 & \text{if } a = b; \\ 0 & \text{otherwise.} \end{cases}$$

we get  $k(A, B) = |A \cap B|$ , the same kernel as before.

# Kernel design (II): practical issues

## Convolution kernels

- Kernels between composite objects by decomposition on the their respective parts, compare components and combine results
- Given two composite objects  $x, x'$  and a way  $\mathcal{D}$  to obtain all possible decompositions into a finite number  $P$  of parts, their  $\mathcal{D}$ -convolution is a kernel:

$$k_{\mathcal{D}}(x, x') = \sum_{x_0 \in \mathcal{D}(x)} \sum_{x'_0 \in \mathcal{D}(x')} \prod_{p=1}^P k_p(x_0(p), x'_0(p))$$



# Kernel design (II): practical issues

## Convolution kernels


Concatenations  
of strings

abracadabra:  
a + bracadabra  
ab + racadabra  
abr + acadabra  
abra + cadabra  
...

Subtrees



Legos (?)



Examples of decomposable objects

# Kernel design (II): practical issues

## P-Kernels

Given a probability distribution on  $\mathcal{X} \times \mathcal{Z}$ , we can compare data points by assigning a high value if both have high conditional probability:

$$k(\mathbf{x}, \mathbf{x}') = \sum_z p(\mathbf{x}|z)p(\mathbf{x}'|z)P(z) \quad \mathcal{Z} \text{ is discrete}$$

$$k(\mathbf{x}, \mathbf{x}') = \int p(\mathbf{x}|z)p(\mathbf{x}'|z)p(z) dz \quad \mathcal{Z} \text{ is continuous}$$

The feature maps are  $(\phi(\mathbf{x}))_z = p(\mathbf{x}|z)\sqrt{p(z)}$

---

Idea:  $p(\mathbf{x}, \mathbf{x}', z) = p(\mathbf{x}, \mathbf{x}'|z)p(z) = p(\mathbf{x}|z)p(\mathbf{x}'|z)p(z)$

# Kernel design (II): practical issues

## Marginalized kernels

Given a probability distribution on  $\mathcal{X} \times \mathcal{Z}$ , and a kernel on  $\mathcal{X} \times \mathcal{Z}$  pairs, we can define:

$$k(x, x') = \sum_z \sum_{z'} k((x, z), (x', z')) p(z|x) p(z'|x')$$

Typical applications of both kernels are found on **graphical models**:

- $\mathcal{X}$  are the *observed* variables and  $\mathcal{Z}$  are the *hidden* variables
- A kernel for the observed ones is obtained by taking the expectation w.r.t. the hidden ones (*marginalizing* them away)
- Examples: HMMs for sequences or stochastic context-free grammars for RNA sequences –see “Kernel methods in genomics and computational biology” (J.P. Vert)

# Kernel design (II): practical issues

---

## The Spectrum (aka $n$ -Gram) kernel

- Let  $\Sigma$  be a finite alphabet: an  $n$ -Gram is a block of  $n$  adjacent characters in  $\Sigma$

$$\text{Define } k(\mathbf{x}, \mathbf{x}') := \sum_{s \in \Sigma^n} |s \in \mathbf{x}| \cdot |s \in \mathbf{x}'| = \sum_{s \in \Sigma^n} f_s(\mathbf{x}) f_s(\mathbf{x}')$$

---

**Example:** Word aababc in alphabet  $\Sigma = \{a, b, c\}$ ,  $n = 2$ :

aa	ab	ac	ba	bb	bc	...
1	2	0	1	0	1	...

---

- While the feature space is large, the feature vectors are sparse; this kernel can be computed in  $O(|\mathbf{x}| + |\mathbf{x}'|)$  time and memory (the actual number of distinct  $n$ -Grams in a text is very small)

# Kernel design (II): practical issues

## Kernels from graphs

- Consider a graph  $G = (V, E)$ , where the set of vertices (nodes)  $V$  are the data points and  $E$  is the set of edges. Call  $N = |V|$ , the number of nodes
- The idea is to compute a (base) matrix  $B_{N \times N}$  whose entries are the weights of the edges and consider  $B^2 = BB$  ( $B$  need not be symmetric)
- Typical use: **connectivity matrix** of  $G$ : the  $(i, j)$  element of  $B^2$  is the number of paths of length exactly 2 between  $i$  and  $j$

---

Examples:

1. protein-protein interactions
2. people-to-people interactions

---

In 2, the  $(i, j)$  element of  $B^2$  is the number of common friends between data points  $i$  and  $j$  (it can be thought of as a measure of their similarity)

# Kernel design (II): practical issues

## Kernels from graphs

Notes:

1. The entries of  $B$  may be real-valued numbers (e.g., symmetric bounded similarities)
2. Higher powers of  $B$  measure higher-order similarities
3. Only the even powers are guaranteed to be PSD

Consider, for a given  $\lambda \in (0, 1)$ :

$$\sum_{k=0}^{\infty} \frac{1}{k!} \lambda^k B^k = \exp(\lambda B)$$

If  $B = U\Lambda U^T$  is the spectral decomposition of  $B$ , then  $B^2 = (U\Lambda U^T)(U\Lambda U^T) = U\Lambda^2 U^T$ .

In general, we have  $B^k = U\Lambda^k U^T$  and therefore:

$$K = \exp(\lambda B) = U \exp(\lambda \Lambda) U^T$$

is an example of a **diffusion** kernel (the name comes from the *heat equation* in physics)

# Kernel design (II): practical issues

## A worked example

Suppose we have designed a “kernel” on **image** objects  $\boldsymbol{x}, \boldsymbol{x}'$  in  $\mathcal{X} = [-1, +1]^d$  from a set of  $D$  descriptors  $f_i : \mathcal{X} \rightarrow \mathbb{R}$  as the function:

$$k(\boldsymbol{x}, \boldsymbol{x}') = \frac{\left( \sum_{i=1}^D f_i(\boldsymbol{x}) f_i(\boldsymbol{x}') + \theta \right)^3}{\sqrt{\theta - \exp(-\|\boldsymbol{x} - \boldsymbol{x}'\|^2)}}$$

where  $\theta \in \mathbb{R}$  is a free parameter.

Under what conditions does this expression define a kernel?

# Kernel design (II): practical issues

## A worked example

1. The numerator is simply the polynomial kernel of degree 3 applied to the descriptors and hence a kernel in itself for all  $\theta \geq 0$
2. Since the denominator must be well-defined, we conclude that  $\theta > 1$
3. Now define  $f(z) = \frac{1}{\sqrt{\theta - z}}$ , for  $\theta > 1$
4. We find the Taylor series expansion for  $f$  as  $f(z) = \sum_{n=0}^{\infty} a_n z^n$ , with  $a_n = \left(2^n \theta^{(2n+1)/2}\right)^{-1} \prod_{i=1}^n (2i - 1)$ , which is positive for all  $\theta > 0$
5. Therefore both the denominator and numerator are kernels and so is their product, for  $\theta > 1$  (note the kernel is *not* normalized)



# Kernel design (II): practical issues

## Kernels as similarity measures

- A kernel can be conceptually regarded as a similarity measure on the data, though many kernels do not fulfill the classical properties for a similarity (e.g. boundedness)
- In a very general sense, a similarity  $s : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$  should at least satisfy the **Similarity Principle**:

$$s(x, x) > s(x, y) \geq 0, \quad \forall x \neq y \in \mathcal{X}$$

- It seems natural to base kernel design on classical similarity functions, for example on **distance-based** similarities
- In many domains, data is described by a **mixture** of variable types (binary, categorical, real, ...)

# Kernel design (II): practical issues

## Gower-like similarity measures

- For any two vectors  $\mathbf{x}_i, \mathbf{x}_j \in \mathcal{X}$  to be compared on the basis of feature  $k$ , a score  $s_{ijk}$  is defined based on their type
- Set  $\delta_{ijk} = 0$  when the comparison of  $\mathbf{x}_i, \mathbf{x}_j$  cannot be performed on the basis of feature  $k$  (e.g. presence of missing values), or we want to ignore it; otherwise, let  $\delta_{ijk} = 1$

Gower's **similarity** is defined as the average over all partial comparisons:

$$s_{ij} := s(\mathbf{x}_i, \mathbf{x}_j) = \frac{\sum_{k=1}^d s_{ijk} \delta_{ijk}}{\sum_{k=1}^d \delta_{ijk}}$$

## Kernel design (II): practical issues

### Gower-like similarity measures

**Binary** (dichotomous) variables indicate the presence/absence of a trait, marked by the symbols  $+$  and  $-$ . Their similarities are given by:

	Values of feature $k$			
Object $x_i$	$+$	$+$	$-$	$-$
Object $x_j$	$+$	$-$	$+$	$-$
$s_{ijk}$	1	0	0	0
$\delta_{ijk}$	1	1	1	0

# Kernel design (II): practical issues

## Gower-like similarity measures

- A **categorical** variable takes discrete and unordered values, which are commonly known as **modalities**
- No order should be defined (the only meaningful relation being equality / non-equality)

Their *overlap* is:

$$s_{ijk} = \mathbb{I}_{\{x_{ik}=x_{jk}\}} = \begin{cases} 1, & \text{if } x_{ik} = x_{jk}; \\ 0, & \text{if } x_{ik} \neq x_{jk} \end{cases}$$

## Kernel design (II): practical issues

### Gower-like similarity measures

**Real-valued** variables are compared with the standard metric in  $\mathbb{R}$ :

$$s_{ijk} = 1 - \frac{|x_{ik} - x_{jk}|}{R_k},$$

where  $R_k$  is the *range* of feature  $k$  (the difference between the maximum and minimum values)

## Kernel design (II): practical issues

### Gower-like similarity measures

**Theorem.** If there are no missing values, the similarity matrix  $S = (s_{ij})$  is PSD (Gower, 1971).

Consider three objects in  $[1, 5] \subset \mathbb{R}^4$ , that is,  $R_k = 4$  given by:

Feature no.	#1	#2	#3	#4
Object $x_1$	1.0	2.0	3.0	1.0
Object $x_2$	1.0	3.0	3.0	?
Object $x_3$	1.0	3.0	3.0	5.0

## Kernel design (II): practical issues

### Gower-like similarity measures

Then we have

$$S = \begin{pmatrix} 1 & \frac{11}{12} & \frac{11}{16} \\ \frac{11}{12} & 1 & 1 \\ \frac{11}{16} & 1 & 1 \end{pmatrix}, \quad \det(S) = -\frac{121}{2304} < 0$$

and therefore  $S$  is not PSD. However, if we replace ? by *any* value in  $[1, 5]$ , then the matrix  $S$  is certainly PSD.

# Examples in real application domains

## Classification of DNA sequences

- DNA sequences of promoter and non-promoters
- A *promoter* is a region of DNA that initiates or facilitates transcription of a particular gene
- The dataset consists of 106 observations and 57 categorical variables describing the DNA sequence, represented as the nucleotide at each position: [A] adenine, [C] cytosine, [G] guanine, [T] thymine
- The response is a binary class: “+” for a promoter gene and “−” for a non-promoter gene



# Examples in real application domains

## Classification of DNA sequences

**Overlap kernel:** the similarity between two multivariate categorical vectors is the fraction of the number of variables in which they match:

$$k_{\text{O}}(\mathbf{x}, \mathbf{x}') = \frac{1}{d} \sum_{l=1}^d \mathbb{I}_{\{x_l = x'_l\}}$$

Another kernel that can be used is the Gaussian RBF kernel:

$$k_{\text{RBF}}(\mathbf{x}, \mathbf{x}') = \exp\left(-\gamma \|\mathbf{x} - \mathbf{x}'\|^2\right), \gamma > 0$$

In order to use this kernel, categorical variables with  $m$  modalities are coded using a binary expansion representation.

# Examples in real application domains

## Classification of DNA sequences

Univariate kernel  $k_1^{(U)}$ :

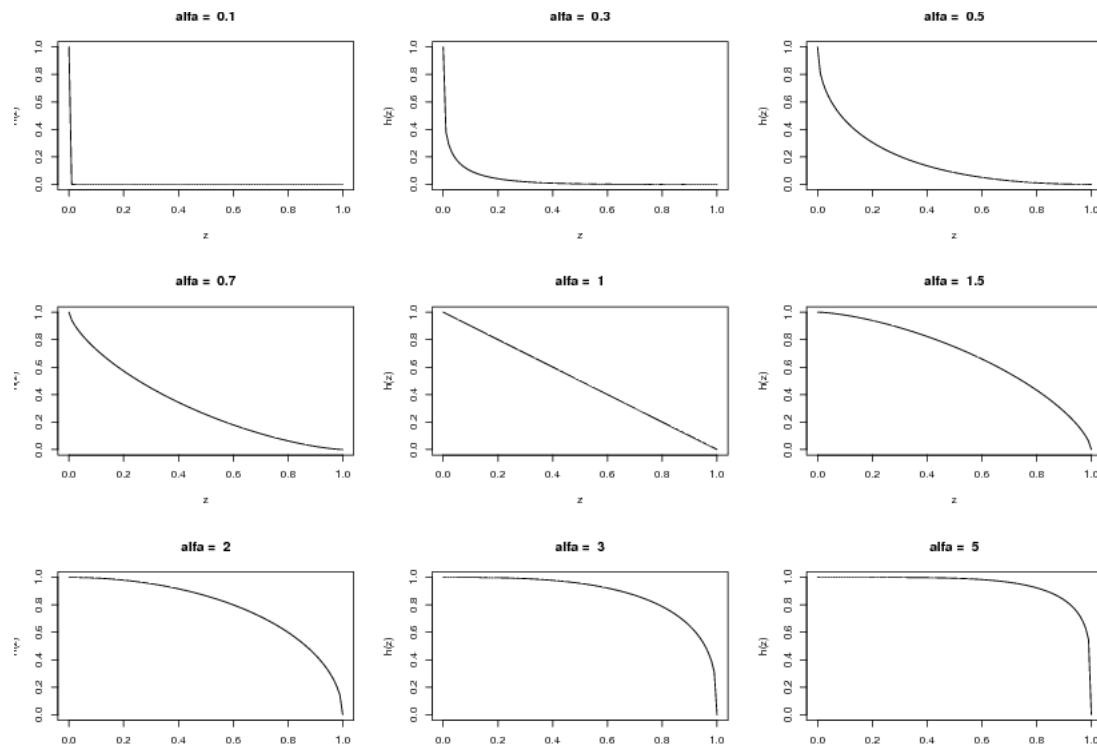
$$k_1^{(U)}(z_i, z_j) = \begin{cases} h_\alpha(P_Z(z_i)) & \text{if } z_i = z_j \\ 0 & \text{if } z_i \neq z_j \end{cases}$$

Inverting function:

$$h_\alpha(z) = (1 - z^\alpha)^{1/\alpha}, \quad \alpha > 0$$

# Examples in real application domains

## Classification of DNA sequences



The family of inverting functions  $h_\alpha(z)$  for different values of  $\alpha$

# Examples in real application domains

## Classification of DNA sequences


Multivariate kernel  $k_1$ :

$$k_1(\mathbf{x}, \mathbf{x}') = \exp \left( \frac{\gamma}{d} \sum_{l=1}^d k_1^{(U)}(x_l, x'_l) \right), \quad \gamma > 0$$

**Proposition.** The kernel matrix generated by the multivariate kernel  $k_1$  is PSD.

# Examples in real application domains

## Classification of DNA sequences



Test error distributions on the PromoterGene problem

*(joint work with M. Villegas)*

## Kernel design (II): practical issues

### More Kernels!

Kernels abound in computational biology and computational chemistry (e.g., phylogenetic profiles, protein 3D structures)

Example: the prediction of **interacting proteins** to reconstruct an interaction network can be posed as a binary classification problem: given a pair of proteins, do they interact or not?

→ we need kernel between *pairs* of proteins!

# Kernel design (II): practical issues

## More Kernels!

The available data is about each single protein; it is then natural to derive kernels for pairs of proteins  $k_{\text{pair}}$  from any kernel  $k$  for single proteins:

$$k_{\text{pair}}((A, B), (C, D)) := k(A, C)k(B, D) + k(A, D)k(B, C)$$

(there is usually no order in a protein pair, so we try both matches)

---

Using Product Kernels to Predict Protein Interactions. Advances in Biochemical Engineering/Biotechnology (110), pp 215-245 (2007)

Kernel methods for predicting protein-protein interactions. Bioinformatics. 2005

# Kernel design (II): practical issues

## Conclusions

- The power of kernel methods partly relies in the ability to process virtually any sort of data as soon as a valid kernel is defined
- Importance of designing kernels that do not constitute explicit inner products between objects, and therefore fully exploit the kernel trick
- Possibility of learning the kernel function (or the kernel matrix) from the training data
- Theoretical analyses are needed on the implications of the kernel choice for the success of specific kernel-based methods